

Student Data & AI Tools

What You Need to Know to Protect Privacy

Emily Beam

Thinking with Agents · Teaching Webinar · June 17, 2026

Why Student Data Protection Matters

"I'd love to use AI for [grading / feedback / course prep], but I'm worried about student privacy."

This concern is legitimate — and universal. Regardless of jurisdiction:

- ▶ Students trust us with **sensitive information**: grades, work, personal details
- ▶ AI tools **transmit data to cloud servers** — that changes the risk profile
- ▶ Data protection law is evolving fast, but the **ethical obligation** is constant

Privacy rules are **not a reason to avoid AI**. They're a reason to **use it carefully**.
The practical protections are the same whether you're bound by FERPA (US), GDPR (EU), PIPEDA (Canada), or institutional policy alone.

FERPA: One Concrete Framework

For US institutions: FERPA (20 U.S.C. § 1232g) is the primary federal student privacy law.



Core Protections

Education records = records directly related to a student, maintained by the institution

PII = names, IDs, grades, and anything that could identify a student alone or in combination

Key nuance: personal notes lose their FERPA exemption the moment they enter a cloud AI tool (no longer “sole possession”)



When Sharing Is Allowed

School official exception: third-party vendors can access records IF the institution has a contract, maintains direct control, and limits use to the stated purpose (34 CFR § 99.31(a)(1))

De-identified data: when all PII is removed and no “reasonable basis” for re-identification

50 years, zero fines — but institutional policy violations and state laws are real risks

Non-US? GDPR, PIPEDA, and national data protection laws impose *stronger* requirements. The practical guidance in this segment applies everywhere — FERPA is the floor, not the ceiling.

Outside the US: What Else Is Required?

FERPA is relatively permissive. Other frameworks add obligations your slides should reflect.

Obligation	US (FERPA)	EU (GDPR)	Canada	Australia
Privacy Impact Assessment before use	Not required	Mandatory	Mandatory	Recommended
Tell students AI is processing their data	Not required	Required	Required	Required
Human review of AI-assisted grading	Not required	Required	Not explicit	Not explicit
Cross-border transfer mechanism	N/A	SCCs or DPF	PIA + notice	APP 8
Student right to erasure	None	Required	Limited	Correction only

Three things non-US faculty should do that US faculty often don't:

1. **Tell students** in your syllabus that AI tools process their work.
2. **Ask your privacy office** whether a DPIA/PIA has been completed for the tool you want to use.
3. **Never let AI be the sole decision-maker** for grades or academic placement.

The Key Question for Any Jurisdiction



Does Your Institution Have an Agreement with the AI Vendor?

Under FERPA, GDPR, and most institutional policies, a third-party tool can handle student data **if the institution has contractual control**. The contract must ensure:

1. The institution has **direct control** over how the vendor uses student data
2. Data is used **only for the stated purpose** — no model training, no secondary use
3. The vendor does not **re-disclose** student data to third parties

✓ Institutional tools

Enterprise agreements with your university (e.g., Microsoft Copilot with DPA, institutionally licensed Gemini, Brightspace)

✗ Consumer tools

No institutional agreement, no control over data use (personal ChatGPT, free Copilot, Claude Free)

“Not Trained On” $\neq$ “Stays On Your Machine”

The most common misconception about AI tool privacy:



What Actually Happens When You Use an AI Tool

1. You open a file → the AI reads it into context
2. File contents are **transmitted** to the provider's servers
3. The provider processes it, generates a response, sends it back
4. The provider **retains your data** (typically 30 days, even on paid plans)
5. A copy is saved **locally** as an unencrypted transcript

“Not used for training” = provider won't put your data into future model weights.

It does **not** mean your data never leaves your machine.

If your obligation requires data *not be transmitted to third-party servers*, no-training is **insufficient**.

Enterprise vs. Consumer Tools

Whether an AI tool is appropriate for student data depends on the **plan tier**, not the tool name.

	Inst. agreement?	Trained on your data?	Retention
UVM Copilot (M365 A5)	✓ Yes	No	Per MS DPA
ChatGPT Free/Plus	✗ No	Yes, by default	30 days
ChatGPT Enterprise/Edu	✓ If contracted	No	Configurable
Claude Free/Pro	✗ No	Opt-in (off default)	30 days
Copilot Free	✗ No	Yes, default (Apr '26)	2 years

Rule of thumb: If your institution has an enterprise agreement with the provider, you have the same protections as any other university-provisioned software.
If not, treat it like emailing student data to a stranger.

Four Levels of Protection



Level 1: De-identify first

Easiest

Strip names, student IDs, and emails before the AI sees the data.



Level 2: Use institutional tools

Recommended

Enterprise tools with institutional agreements create a FERPA-compliant channel.



Level 3: Sandbox the AI tool

Technical

Restrict which files and directories the AI can access. Kernel-enforced.



Level 4: Physical separation

Strongest

Keep student data on a separate drive or encrypted volume. Unmount before AI sessions.

Stronger
↓

Most faculty tasks are fully addressed by **Levels 1–2**. Combine levels for defense in depth.



De-identification is a wood chipper for PII.

Identifiable data goes in. Anonymous content comes out.

De-Identification in Practice

✕ Before (contains PII)

Name	ID	Grade
Jane Smith	901234	87
Alex Chen	905678	72
Maria Lopez	903456	91

✓ After (safe to share)

Student	Grade
Student_001	87
Student_002	72
Student_003	91

⌕ A simple Stata script handles this

```
use "grades.dta", clear
drop name student_id email
gen student = "Student_" + string(_n, "%03.0f")
order student
export delimited using "grades_anon.csv", replace
```

For most assignments (papers, problem sets, reading responses), stripping names and IDs is sufficient. Exercise more caution with personal narratives or very small classes where content itself could identify a student.

The Distinction That Matters

System-to-System (identity attached)

Student submits through LMS → AI tool receives the paper **with name, course, and metadata**

Examples: Turnitin AI detection, LMS analytics plugins, Copilot pulling from your SIS

✗ **Requires** institutional contract or student consent.

Manual De-Identification (link broken)

You copy a student's answer, strip the name, paste anonymous text into an AI tool

Examples: Problem set for feedback help, reading response for comment drafting, anonymous essay for pattern analysis

✓ **Fine for most assignments.** You broke the link.

Privacy law protects the link between a student and their work — not the academic content itself.

Break the link, and the content is no longer a protected education record.

The “Helpful AI” Problem

Even with safeguards, AI tools optimize for **helpfulness** — and that creates risk.



Why Defense in Depth Matters

- ▶ AI agents can **read any file** your user account can access — by default
- ▶ They will proactively seek out context to give you a better answer
- ▶ A roster, gradebook, or email export in the same directory? **Fair game.**
- ▶ Permission prompts are **not a security boundary** — they can be bypassed by bugs, subprocesses, or misconfigurations

Don't rely on a single guardrail. Combine protections:

De-identify the data **AND** use an institutional tool **AND** keep student files out of the working directory.

“Can’t I Just Run a Local Model?”

You may hear that **local models** — AI that runs entirely on your laptop — solve the privacy problem. In theory, yes: no data leaves your machine. In practice, this is not a realistic option for most faculty.

! Why local models aren’t the answer (yet)

- ▶ **Significantly less capable** than cloud models (GPT-4, Claude) — worse at nuanced feedback, rubric application, and writing
- ▶ Requires **technical setup** that most faculty don’t have time for (Ollama, LM Studio, command line)
- ▶ Needs a **powerful computer** — 16GB+ RAM for usable models, and they run slowly
- ▶ **No institutional support** — if something breaks, you’re on your own
- ▶ The quality gap means you’d spend more time **fixing AI output** than you saved

The simpler path: De-identify your data (strip names) and use a capable cloud tool. This solves the privacy problem without sacrificing quality or requiring technical expertise.

Common Scenarios

I want to...	Risk?	What to do
Generate feedback templates from my rubric	None	No student data involved. Use any tool.
Use AI to draft comments on student papers	Medium	De-identify first, or use institutional tool.
Analyze grade distributions	Low	Strip names/IDs. Aggregated stats are fine.
Build a course website or syllabus	None	No student data involved. Use any tool.
Create an AI chatbot for students	High	Student interaction data may be education records. Consult IT.

What to Tell Your Department Chair

Template Language for Department Policies


“Faculty may use AI tools for course preparation and instructional support. When student data is involved:

1. **Use institutionally approved tools** with enterprise/data processing agreements.
2. **De-identify student data** before using any AI tool not covered by an institutional agreement.
3. **Never paste identifiable student records** into consumer AI tools.
4. **Keep student data files separate** from directories where AI tools operate.

These practices are consistent with FERPA (US), GDPR (EU), and most institutional data governance frameworks.”

Adapt for your institution. IT security or your data protection officer can confirm which tools are covered.

Quick Decision Guide for Teaching Data



Key Takeaways

1. **Privacy law is not a ban on AI.** FERPA, GDPR, and institutional policies regulate *how* student data is handled — and provide clear paths forward.
2. **Institutional tools are the safe path.** Enterprise agreements give your university contractual control over data use.
3. **De-identification is your simplest tool.** Strip names and IDs before the AI sees the data. But beware small classes.
4. **“Not trained on” is not enough.** Your data still leaves your machine. Understand what the tool actually does.
5. **Defense in depth.** Combine de-identification + institutional tools + file separation.

Resources



Key Resources

- ▶ **FERPA:** 20 U.S.C. § 1232g; 34 CFR Part 99 (§ 99.31, § 99.3) · **GDPR:** Articles 6, 22, 35
- ▶ **Future of Privacy Forum:** Checklist for vetting AI tools (2024) — fpf.org
- ▶ **Oregon State AI Decision Tree:** Faculty flowchart, CC-BY-NC — ecampus.oregonstate.edu
- ▶ **AAUP:** “AI and Academic Professions” (2025) — survey of 200 campuses on governance gaps
- ▶ **UNESCO:** “Guidance for GenAI in Education and Research” (2023) — first global framework
- ▶ **CDT State Legislative Tracker** (2025): 53 bills / 21 US states — cdt.org

Slides, de-identification scripts, and the full data protection brief
available at thinkingwithagents.github.io